This week, we’re having a test. It is available in Edux, in the “Tests” section.

Lecture 11: Motherboard

ROM and booting

As we know from the previous lectures, the role of RAM includes storing the code of the currently running program. However, at the time when a computer is powered on, RAM does not contain any useful data. The operating system, which also participates in memory management, isn't yet active either. To make setting up all of this possible, the computer motherboard contains an additional memory die. Clearly, it must be **non-volatile memory**, i.e. one which does not lose data during power-off periods.

That memory is usually of the **flash EEPROM** type. The name *EEPROM* can be a bit confusing, as it stands for: *electrically erasable programmable read-only memory*. The *read-only-memory* (*ROM*) part is here for historical reasons. Indeed, several production technologies ago, an important spot in computer design was occupied by memory in which data could be stored only once, at manufacturing time. However, later enhancements enabled that memory to be edited ("programmed") long after its creation (leading to PROM), then to be erased with UV light and programmed again (EPROM), up to the modern solutions in which such memory can be erased and programmed again without even taking it out of the computer.

The main content of the built-in non-volatile memory is *firmware* purposed to be executed as the first code after the computer is powered on. This plays the role of a "restricted operating system", enabling the processor to communicate with input/output devices on the motherboard (e.g. keyboard, disk, graphics card) until the true (much more complex) operating system is started and takes these tasks on.

For a long time, that initial firmware was usually based on a program for the IBM PC computer called BIOS — and, despite later extensions and adjustments, used to be still generally described as **BIOS** (*Basic Input/Output System*). In fact, that name can be still found in use, e.g. the abovementioned ROM memory is still sometimes called the *BIOS die*. However, in the last years, due to technical limitations typical to BIOS, its role has been taken over by a new firmware standard called **UEFI** (*Unified Extensible Firmware Interface*), developed and promoted in cooperation between multiple leading companies (including Intel, AMD, Lenovo, Dell, HP, Microsoft, and Apple).

Before going into the differences between BIOS and UEFI, we will describe the basic principles of the startup procedure of a computer and its operating system, which look similarly in both standards.

Booting

One of the main tasks of BIOS/UEFI is executing a **boot loader**, responsible for loading the operating system code into RAM, and then executing that code. In modern computers, this procedure often consists of **multiple stages**:

- A *first-stage* boot loader (BIOS/UEFI), instead of loading the operating system, loads another *second-stage* boot loader, which can be e.g.:
  - a boot loader for MS Windows (`bootmgr`), offering the user a choice between Windows versions (if there is more than one installed), as well as an option to run Windows in so-called *safe mode* (a restricted, and hence more stable, version of the system);

  - the GRUB boot loader, offering the user a choice between various UNIX-like operating systems (including Linux) and MS Windows — of course, if there is more than one such system installed.

- Only then, the proper operating system is invoked.

The second-stage boot loader is initially placed in the secondary memory, in a well-specified location of the particular storage device (e.g. for a hard drive, it is its first sector). (A part of a first-stage boot loader's work is to check all the computer storage drives for one that has such a boot sector). One can also boot from a pendrive, a CD/DVD disc, or even from another computer in the local network. (Earlier in the past, there were also other technologies in use, like floppy discs, or even punched tapes or cards). Notably, during an installation procedure for an operating system, the second-stage boot loader role is taken by the installer.

The POST self-check

Before BIOS/UEFI can proceed to the boot loader, it executes the **POST** procedure, which is the basic check of the crucial computer components. In particular, the checks involve the processor and main memory. In case any erorrs have been found, the appropriate message will be displayed on the screen — however, in case this could be impossible (e.g. due to a screen failure), the problems are additionally communicated with sound beeps. The exact meaning of these signals depends on the BIOS/UEFI manufacturer. Below, we list the meanings of sound signals according to one of the popular conventions, AMI BIOS:

Signal	Error type	Comment
1 long	no error	the POST test has finished with success
1 short	DRAM refresh error	e.g. an error while handling an interrupt
2 short	memory parity error	unexpected value of the checksum bit
3 short	RAM error	within the initial 64 kiB
4 short	system clock error	
5 short	CPU error	
6 short	A20 gate error	the *A20 gate* controls CPU access to memory addresses above 16 MiB; an error in it prevents CPU from switching between modes of operation
7 short	CPU virtual mode error	
8 short	video memory error	
9 short	ROM checksum error	ROM content has been corrupted
10 short	CMOS memory error	*CMOS* is a small memory for BIOS settings (which is non-volatile thanks to permanent power supply from the *CMOS battery*)
11 short	L2 cache error	
1 long, 2 short	video subsystem error	
1 long, 3 short	RAM error	outside the initial 64 kiB

1 long, 8 short	display error	
alternating high/low	insufficient power	insufficient CPU voltage or fan rotation speed

Occasionally, it's worth to check certain components even more thoroughly, which can be achieved with special diagnostic programs. Memory or hard drive can be tested from the operating system level; however, some tools (often these more thorough, e.g. memtest86 for RAM) need to be launched from the BIOS/UEFI level.

Settings

After POST finishes, the user can enter the settings mode. To land there, one needs to press a proper key — usually Esc or one of the function keys (F8, F12 etc.). The proper key depends on the BIOS/UEFI version; an information regarding which one is proper is displayed for some time on the startup screen.

In the settings mode, the user can adjust e.g. the clock frequency for the CPU, RAM, or system buses, the voltage powering the motherboard slots, or the computer fan rotation speed. Changing these parameters may improve computer efficiency — however, one should keep in mind that this brings a risk of unstable operation, or even damaging the hardware. Fortunately, at least as long as no component has been damaged, the settings mode allows restoring the factory defaults.

In that mode, the user can also set the order in which storage drives are checked for the boot loader. This is practically necessary if the top priority is currently assigned to the hard drive, while we want to boot from a pendrive (theoretically, an alternative is to take the hard drive out of the computer, though in practice it's clearly easier to change the settings). On the other hand, putting the hard drive on top of that list can save a few seconds of waiting during every normal computer startup. Similarly, some time can be saved by switching unnecessary devices — for example, when there is a non-connected slot on the motherboard. Another benefit from such settings change can be increasing security (e.g. by switching off all USB data transmission).

I/O interface

The last functionality of BIOS/UEFI to be mentioned here is defininig interrupt handlers — that is, defining a way for the CPU to communicate with the input/output devices. Until the 1990's, BIOS served as an intermediate layer between the operating system and I/O devices (which is reflected in its name). Currently, managing I/O devices is done inside the operating system; thanks to this, the user can attach a new device (and, if needed, install its software driver) without upgrading their BIOS/UEFI.

However, I/O management in BIOS/UEFI is still retained — e.g. to ensure backward compatibility (so that older programs can still execute), though most importantly to enable the abovementioned booting procedure during which BIOS/UEFI must load data from secondary storage without any help from the operating system (which is not yet running at that time).

BIOS vs. UEFI

From the user's viewpoint, the most noticeable difference between the newer UEFI and the older BIOS is appearance: while BIOS displays everything in ASCII mode, with no mouse support, UEFI offers a graphical interface (and, generally, more convenient usage). UEFI also removes some problems and limitations present in BIOS regarding handling hard drives (e.g. total size limited to 2 TiB, the number of so-called *primary partitions* limited to 4). However, the impact extends to other computer components. Also, some settings (e.g. Secure Boot) can on one hand improve security by making malware attacks significantly harder, but on the other hand may complicate or even prevent installing some operating systems.

UEFI also offers an operating mode compatible with the BIOS standard (*Compatibility Support Mode* — CSM; earlier also *Legacy BIOS*). Enabling it is necessary in case of some older components incompatible with the UEFI standard. On the other hand, some modules (e.g. Intel chipsets) do not offer compatibility with the BIOS standard. Similarly, the MS Windows operating system — starting from version 11 — requires the Secure Boot setting to support some of its functionalities, and that is only available in the UEFI mode, not in the BIOS compatibility modes.

Upgrades

As already mentioned, BIOS/UEFI on a particular computer can be upgraded (which may be necessary e.g. after replacing a hardware component). Saving a new version can be done in various ways (listed below starting from the most high-level ones):

1. From the level of an already running operating system (provided that the OS supports this);
2. From a special file placed on a pendrive, during booting, using the UEFI user interface;
3. From a special file placed on a pendrive, by pressing an appropriate button on the computer case (provided that the motherboard and the computer case both support this);
4. Directly to the EEPROM die (after opening the computer case), by using a dedicated device.

The more high-level ways are generally safer (e.g. the operating system and UEFI make an initial check of the update file, so that we do not risk e.g. using a "neighbouring" file of a wrong format by plain mistake). Still, such an update is potentially dangerous — it could make some I/O device inaccessible. If we're unlucky enough to pick a UEFI version incompatible with our motherboard, we may even lose way to turn the computer on; in such case, the only remaining ways to fix this are methods 3 and 4. The takeaway is that upgrading BIOS/UEFI without a direct reason is not recommended.

Motherboard

The computer parts discussed so far — CPU, RAM dies, the BIOS die — are physically located on the **motherboard**. It is, as the name suggests, the place for attaching the most important components. Besides just hosting them, it enables them to communicate with each other.

On a modern motherboard, we will find e.g. those elements:

- a CPU slot (or *socket*) — a place to put a CPU in
  (though it also happens that the CPU is soldered directly to the motherboard);
- RAM slots: typically DDR3, DDR4, DDR5
  (here, it may also happen that memory is soldered directly to the motherboard[1]);
- the chipset (which will be discussed below);
- the BIOS die;
- the system clock generator, allowing synchronization between various components;
- power sockets, through which all the computer components are powered with electricity;
- hard drive slots;
- optical drive (CD/DVD) slots;
- keyboard slots;
- extension card slots
  (it is through them that we can connect a graphics card or a network card, or a USB port controller).

In practice, a motherboard can look as follows:

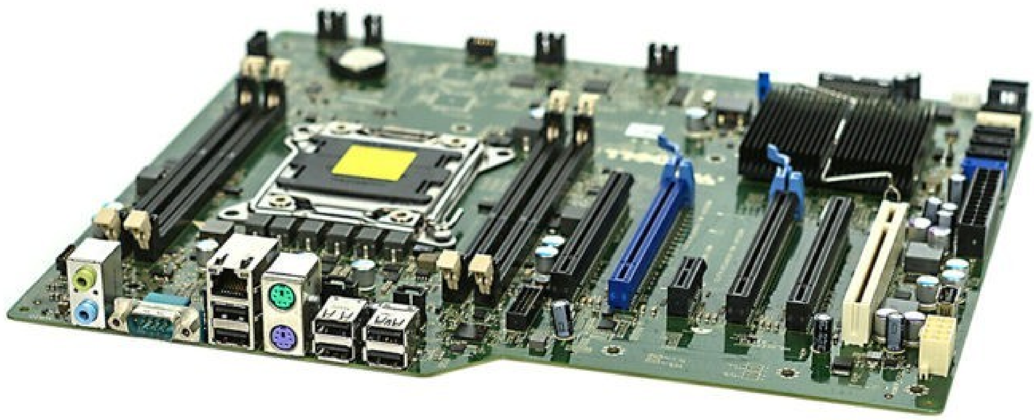

Figure 1. A Dell Precision T3600 motherboard (from year 2012).
(Source: `http://en.wikipedia.org`)

Such organization of the board with standardized sockets/slots allows for some elasticity in choosing the components to be placed inside or attached to the computer. Unfortunately, even if a component can be attached physically, it's still not guaranteed that it will be supported correctly by the motherboard and the processor. Hence, it's essential to verify compatibility specifications on every hardware upgrade.

[1]This happens e.g. for the LPDDR memory, which is used mainly in mobile devices (phones, tablets), though also in certain laptops. This memory kind is slower and more expensive than typical DDR but in reward it requires lower voltage, which allows longer operation on battery power.

Chipset

A **chipset** is literally "a set of (co-operating) integrated circuits". In the case of the motherboard chipset, we can distinguish (or rather "used to distinguish", though the recent changes will be discussed later) two such circuits, called the **northbridge** and the **southbridge**. These are the chips implementing the connection between the processor and the I/O devices. A scheme of the two-bridge architecture looks as follows:

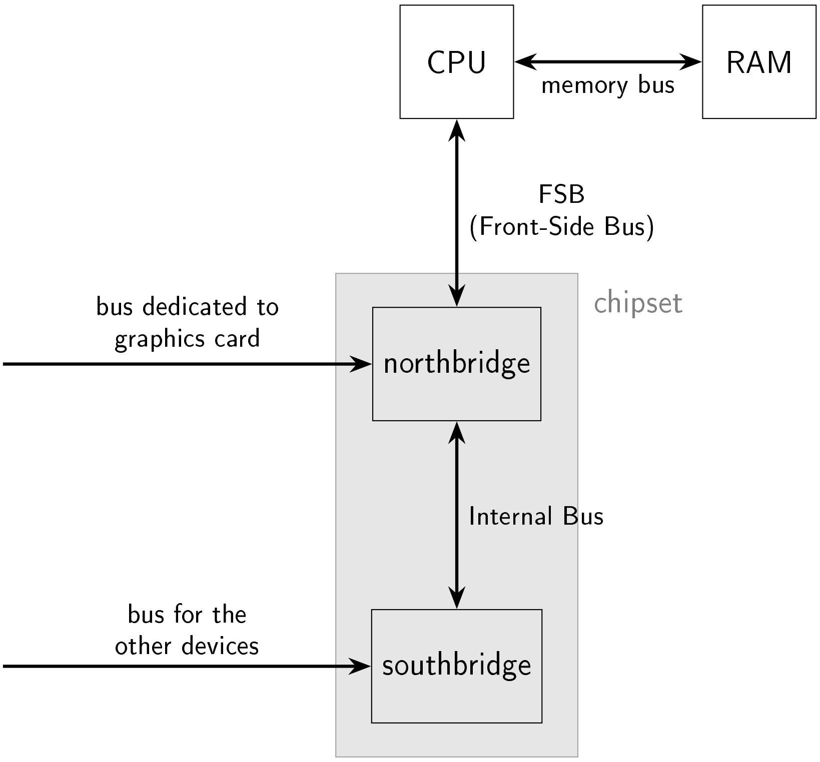

Figure 2. A motherboard chipset with the two-bridge architecture.

As we see, the northbridge is located physically closer to the CPU. This is important, as it connects the components requiring quickest access to the memory. In the early motherboards with the two-bridge structure (e.g. in the CS8220 chipset, which introduced a component called "northbridge" in 1986), the northbridge used to also contain a memory controller connected to memory buses. In the more modern solutions, as we mentioned above, these have been integrated into the processor.

The southbridge is not connected to the processor directly, but via the northbridge. That one connects — directly or indirectly — the processor to all the other (slower) input/output circuits.

In the most modern solutions, the tasks served so far by the northbridge have been taken over by the processor. This looks as follows:

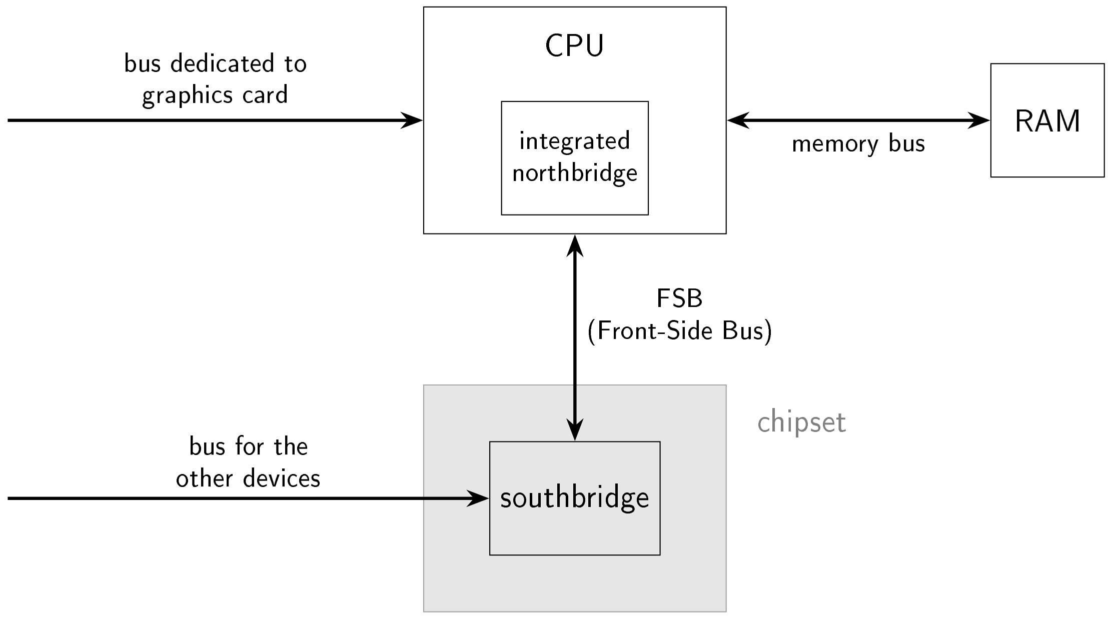

Figure 3. A motherboard chipset with the one-bridge architecture.

Even in a so basic description, we can see that various computer components can be either of:

- a separate integrated circuit,
- integrated with the motherboard,
- integrated with the processor.

This choice is not visible from the user viewpoint, as long as there is no need of replacing some parts. This means that a higher level of integration leads to a risk of more expensive servicing once some components broke, or of hardware replacement becoming more difficult. On the other hand, the advantage of a higher level of integration can be better efficiency.

Altogether, this means that also in the future we can expect more changes of computer architecture in terms of moving various components between the above categories (just as it happened to the northbridge, and somewhat earlier — at least roughly — to the RAM controller).

Buses

Buses play a key role on the above diagrams: they transmit signals between I/O devices, the chipset, and the processor. Their bandwidth capacity and transfer speed is one of the crucial points for the overall computer efficiency.

Currently, the standard for high-speed data busses is PCI Express. It allows transmission in both directions at once (*full duplex*). Transmission takes place on multiple **lines**, which allows sending simultaneously multiple bits in one direction. The newest version offer in total a bandwidth of tens of gigabytes per second. However, not every input/output device needs so high bandwidth; therefore the southbridge is also connected to buses in many other standards. We will touch the topic of buses again when discussing the input/output system.

Lecture 12: Non-volatile storage

So far, we've taken a closer look at the primary memory. Now, it's time to discuss the most important kind of non-volatile memory, which is **mass storage** (including: hard drives, SSD drives, but also external memory carriers such as CD/DVD discs or USB memory). Although theoretically a computer can operate without any memory of such type (e.g. in the mode "receive the program to be executed from the network and print the output on a printer"), this is very rare in practice. A typical (personal or office) computer is usually equipped with a built-in mass storage, popularly called a **drive** (or a *disk*), which should be regarded as one of the basic computer components.

Storage drives

Below, we'll discuss two main kinds of such large devices: the HDD and SSD drives.

To begin with: a note on the language

Before getting into the details, let us make a language-related digression: there is some confusion about the phrases *drive*, *disk*, *hard drive* and *hard disk*. Although they differ in meaning, in common language all of them happen to be used interchangeably.

A **drive** is defined by the Cambridge Dictionary as "a *device* for storing computer information", while a **disk** refers to the proper data container, assuming (at least in theory) that it's flat and round. That's why, back in the 90's, we would insert a *floppy disk* (a data container) into a *floppy (disk) drive* (a device), or store our files on a *hard disk* (a metal plate) inside a *hard disk drive* (a device for storing data there). At that time, however, "*hard disk drive*" — or any sub-phrase of that — became commonly treated as a synonym of "large secondary memory", just because the latter was rarely implemented in any other way.

Later, when the *solid-state drives* built on flash memory appeared, they kept being commonly called *disks*, despite not containing anything of that shape. Years are passing — but still, whenever abstractly speaking about the secondary memory layer in the computer architecture, we may often hear it called "the *disk*" or "the *hard drive*", even though the actual implementation may be SSD-based (and then, strictly speaking, neither *hard* nor a *disk*).

To add another bit of chaos, the optical data containers have been traditionally described with the British orthography: we'll more often see a *DVD disc* than a *DVD disk*.

Hard drives

The **hard disk drives** (**HDD**, also known as *hard drives* or *hard disks*) are a technology of long history (the first such drive was manufactured by IBM in 1956). They're based on an even older idea of **magnetic storage**, in which information is represented in the local magnetic field of the data container. In every cell, we can create a permanent magnetic field in one of two possible ways (directed in one direction or another), which allows encoding binary digits. (However, the actual encoding is a bit more complex than "encode every bit in the magnetic orientation" — instead, it enforces frequent changes of the orientation, so that any larger area of the disk observes a negligible total magnetic field). In order to magnetize the cells, a head element is used that is capable of creating a targeted (and frequently changing) electromagnetic field. The head also allows testing magnetization of the disk, and hence reading the data.

Physical structure

Due to the necessity of placing the moving head above an arbitrary cell of the data storing material, the shape of that material becomes essential. In the HDD drives, these are flat and round carriers, called **platters** (and originally *hard disks*). A single HDD drive may consist of several such platters: typically 1 or 2 in laptops, but up to 4 in desktop computers. The platters keep rotating, typically with a speed of 7,200 rpm (*revolutions per minute*), i.e. 120 revolutions per second, though there still exist cheaper disks with 5,400 rpm. Consecutive data are stored on along concentric arc paths, called **tracks**.

The following figures show a real look of an HDD drive, and a diagram of its most important components:

Figure 1. An internal view of a HDD drive from the late 1990's.
(Source: `http://en.wikipedia.org`)

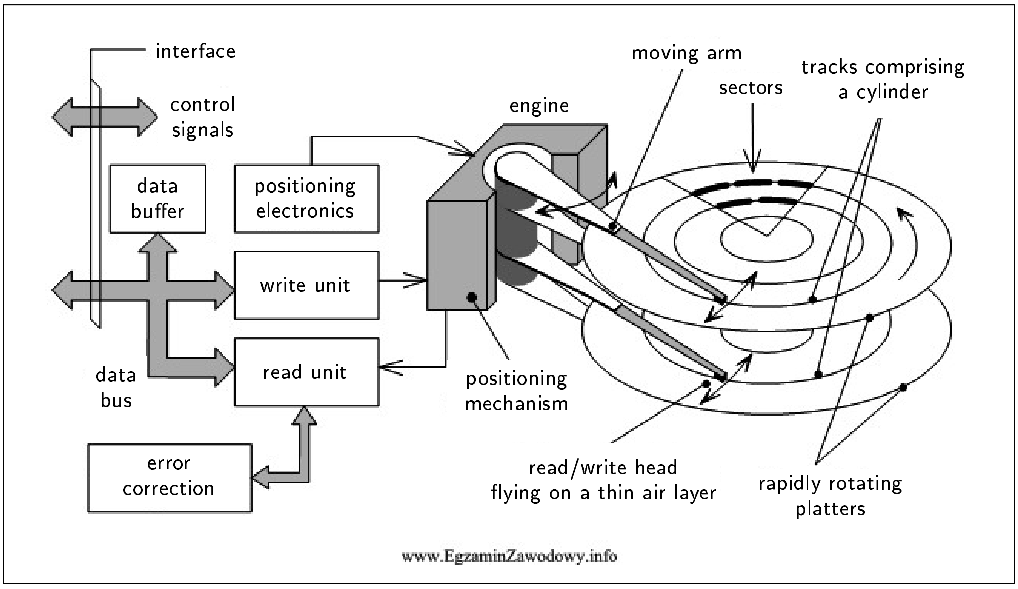

Figure 2. A diagram of a HDD drive structure and its communication with the computer.

The capacity of a hard drive depends on the number of tracks and on the density of information on a single track. The main obstacles for increasing that density are magnetic effects: if single cells are too small and placed too close to each other, they can sometimes flip their magnetic orientation to the opposite one — under the influence of neighbours, or even just by themselves. Until the 2000's, this forced the industry to use the so-called **longitudal recording** (in which the magnetic poles were directed parallelly to the platter surface), which in turn restricted HDD capacity to ca. 1 TB. Currently — thanks to several technological novelties (including a new magnetic material, and a new type of write head) we can employ **perpendicular recording**, which means smaller cells and hence greater capacity.

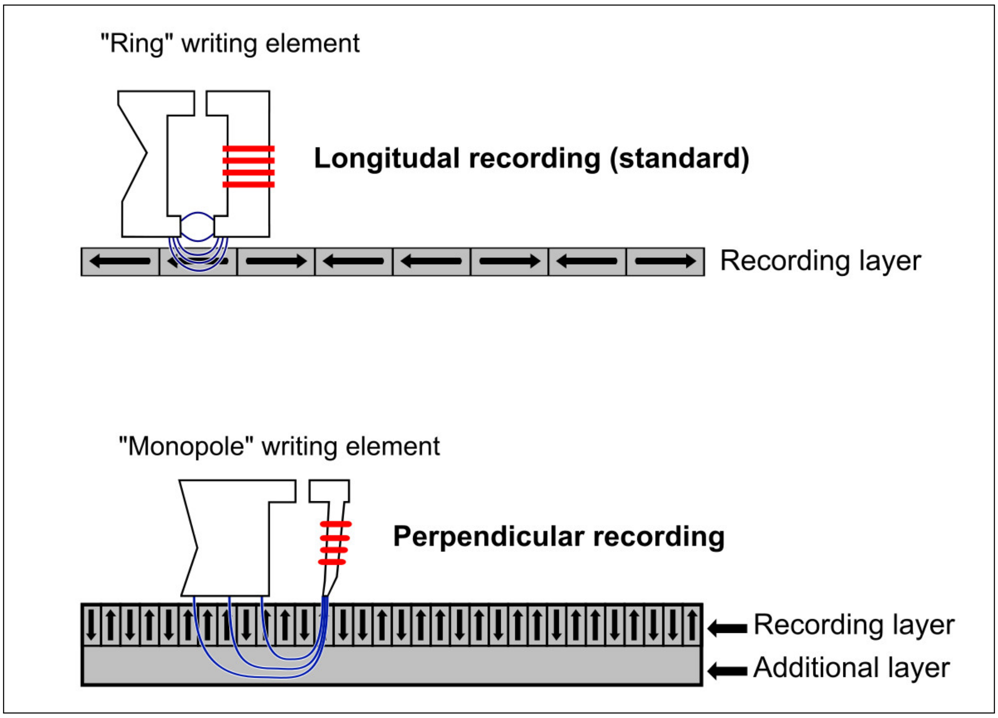

Figure 3. Porównanie zapisu równoległego (u góry) i prostopadłego (u dołu).
(Source: `http://en.wikipedia.org`)

One of the technologies that give hopes for further overcoming the abovementioned barriers is *thermally assisted recroding*, in which the platter surface is temporarily heated up by a laser beam. In 2021, this method allowed Seagate to produce a hard drive of 20 TB capacity — however, it is not yet used commercially. Another promising method (but not yet deployed to the market) is *bit-pattern media*, whose aim is — somewhat roughly — to achieve a decrease of memory cell sizes by ensuring their more regular shapes.

Organization of the data

Due to the structure of the disk (accessing the data with read/write heads placed on a single moving arm), the physical placement of the data becomes important for its access time. The most efficient access will happen when all the relevant data is placed together: ideally on a single track, or at least on a set of nearby tracks (to minimize the arm movement). Also, to minimize the number of read/write operations, data are always read or written in chunks. The smallest such chunk, corresponding to a fragment of a track, is called a **sector**; modern computers use sectors of size 512 B or 4 kB. Every sector also stores additional metadata for controlling the disk operation: these include information on the deviation of the rotational speed, as well as a checksum which allows detecting (and sometimes even fixing) cases of single bit corruption.

The problem of placing related data (e.g. single files) in nearby disk sectors is on one hand crucial for the efficiency of the whole computer, but on the other hand not trivial as the existing files are constantly modified (and, in particular, can keep growing). This can lead to disk **fragmentation**, which can mean either dispersing files throughout distant disk sectors (leading to slower read and write operations) or appearance of unused sectors which are not worth filling with small portions of files (which in turn decreases the effective capacity of a disk). Managing the placement of data on the disk and fighting the fragmentation issue is a responsibility of the software layer — and, specifically, is done by a particular **file system** used by the operating system — which we will describe in more detail below.

SSD drives

As we mentioned before, HDD drives are one of the two main types of internal mass storage that is currently broadly used. The other one are **solid-state drives** (known as *SSD drives*), based on the semiconductor **flash** memory.

The flash memory, proposed in 1980 and manufactured industrially since 1989, is built of special **floating-gate transistors** (*FGT*), equipped with an additional "floating gate" which is surrounded by an isolated material which allows charging it with electric change, with a *long-lasting effect*. (That's a substantial difference in comparison with flip-flops discussed on previous lectures, in which holding any information requires a constant electric powering). This makes FGT transistors suitable as non-volatile memory cells. Still, it's worth realizing that building mass storage out FGTs involves multiple technical complications — for example, in the early versions of this approach, writing to memory was only possible after a prior erasure (called a *flash* operation; hence the name of this technology). Later developments led to the state which we know today, in which SSD drives offer (from the user's perspective) almost unrestricted read and write capabilities.

In terms of the internal structure, SSD drives resemble the already discussed ROM memory dies; in particular, they contain no rotating elements. This allows significantly higher data access times in comparison with the HDD technology. Also, SSD drives are more resistant against heat, physical shocks etc. Their most significant disadvantage is the price, which remains several times higher than for HDD drives.

The SSD drives in home computers typically employ the *multi-level cell* (**MLC**) technology, in which a single memory cell can take more than 2 states: we no longer classify the voltage as just high or low, but instead distinguish between more levels of voltage. These will be typically 4 levels, though it can be also 8 or 16. This allows a single memory cell to store respectively 2, 3 or even 4 bits of information. (Works are in progress to increase that number to 5). Reducing SSD prices to a widely acceptable level was achievable only thanks to using MLC. The disadvantage of this technique is an increased chance for bit errors, and reduction of the lifetime of the drive. Therefore, MLC is not used in BIOS memory (which we discussed in Lecture 10).

External drives

Personal computer users also widely employ **external drives** — which are made with the HDD or SSD technology just as the drives discussed above, but stand out in that they're placed outside the computer case, and can be plugged into it as needed, typically via USB ports. Their obvious advantage is an increase of the total available secondary memory; another clear gain is portability.

Such drives can be also used as data backups for the internal HDD or SSD drives (as both these technologies involve a risk of data corruption or loss) — here, though, one needs to be careful about flash memory: even though essentially labeled as non-volatile, it can lose its content (due to discharging of its internal capacitors) after a few years of not being used.

Other storage means

Let us briefly mention some other input-output devices purposed for storing user data:

- **USB memory** — devices known as *USB sticks* or *pen-drives* (though the latter name is registered by a Taiwanese company Add On Technology, so it officially applies only to their products) — these store data in the semiconductor flash memory, just as SSD drives. The difference here is that they're designed to be external, small and portable. As a result, they're significantly slower (the limit set by the standard is 4.6 Gb/s, though in practice it's usually far from being achieved). Also, they're more error-prone and more expensive (per byte) than SSD drives.
- **SD memory cards** — another type of semiconductor flash memory. Used mainly in photo and video cameras, or in mobile phones.
- **Optical memory** (*CD / DVD / Blu-ray discs*) — here reading and writing is done with a laser beam. The data carrier is a surface (typically of a disk shape) covered with a material whose transparency can be easily changed, thus introducing two types of points (reflecting and non-reflecting) which can represent bit values. Similarly as in HDD drives, reading and writing data employs disk rotation to help in positioning the read/write head above the appropriate sector. The data-carrying materials can support just a single write or multiple writes; they can also differ by total capacity:
  - For the CD standard, this will be typically ca. 1 GB.
  - For DVD, the typical capacity is 4.7 GB; by using double-layered and/or double-sided discs, we can reach up to 17.1 GB, though DVDs of such a large size are rarely used.
  - In the Blu-ray technology, the typical capacity of a disc is 25 GB. Here, even more than two layers can be supported, though it's also a rare case. The increase of information density in this technology is achieved by using a laser beam with a shorter wavelength (a blue light).

Drive interfaces

Besides the manufacturing technology, an important factor for the practical efficiency of a storage drive — and its usefulness in combination with a particular computer — is the way in which it is connected, that is: the type of the bus and the plug used.

Bus types

Currently, in home applications, three main types of buses are used to connect storage drives to the processor:

- **SATA** (*Serial ATA*, where *ATA* stands for *Advanced Technology Attachment*) — practically the only choice for HDD drives; also used for many SSD models;
- **PCIe** (*PCI Express*, where *PCI* stands for *Peripheral Component Interconnect*) — a standard which we already met in Lecture 10 (due to its role in connecting the CPU to the RAM memory); used also for faster SSD models;

- **USB** (*Universal Serial Bus*) — used mostly for attaching external drives.

The SATA bus is derived from an older ATA model, designed specifically for communication with hard drives. For example, its cooperation with the drive controller enables the so-called *native command queuing*, which allows altering the order of executed read and write orders to minimize their total execution time. (Recall that this time depends on the physical placement of data on the disk, so it's beneficial to rearrange the commands so that those involving neighbouring sectors are executed next to each other). This makes SATA a proper choice for HDD drives; this in turn makes it beneficial for SSD manufacturers to support that standard as well, as it allows the end users to replace an older HDD with a newer SSD in a typical computer.

On the other hand, limitations for transmission speed in SATA standard (3 Gb/s in SATA 2 and 6 Gb/s in SATA 3) cause it to be too slow for the fastest currently available SSD drives. Therefore, full utilization of such devices requires connecting them to a faster PCIe bus (where the analogous limit in the fastest version is nearly 1 Tb/s as of 2022, and doubles every few years).

USB stands out from the other two standards in that it's been designed from the start for external devices (and hence we will find USB ports on the outside of a computer case).

It's also worth mentioning here that all the three standards involve, besides transmitting data, supplying the connected devices with an appropriate amount of electric power.

Communication mode

All the above bus types are described as **serial**, which *essentially* means that transmission happens sequentially, bit by bit.

It could seem that a relatively easy and effective idea for increasing the transmission rate could be switching this approach to a **parallel** one, in which multiple bits would be sent at once (or more precisely: in each bus cycle, 1 bit would be sent along every wire). In practice, however, since around 2000 we observe transitioning from older parallel solutions (PCI, ATA) to serial ones (USB, SATA, PCIe). This has happened because mutual interference between the wires, as well as slight differences in transmission times along them, complicate synchronization on the receiving side so much that it led to essential limitations on the frequency of cycles in parallel buses. By migrating to a serial transmission, we achieved significantly higher frequencies.

On the other hand, in a contrast to the word “serial”, modern buses also consist of several wires.

In the SATA standard — which is simplest in this regard — we have 7 wires, out of which 3 transmit no data (serving just as isolation layers for the other ones), and the remaining 4 wires are split into 2 pairs, each responsible for transmission in one direction. In each pair, one of the wires is a strict “negation” the other one, always transmitting the opposite bit values. (In exchange for such “waste” of bandwidth, we gain: a more clear signal interpretation on the receiving side, a chance to detect errors in communication, as well as a constant level of the electric charge on both sides of the bus).

The PCIe buses, as already mentioned in Lecture 10, go much further, introducing **multiple lanes**, each of which allows transmitting 1 bit in each direction. (In practice, each such lane consists of 4 wires, organized into two negation pairs, similarly as in SATA). In version 6.0, using 16 such lanes of transmission rate 60 Gb/s each allows achieving nearly 1 Tb/s in total. (A similar solution — introducing 2 lanes of communication — has been also introduced in the newest versions of USB).

If so, one could ask: why is PCIe still called *serial*? The reason, at least partially, is that each PCIe lane operates as an independent serial connection, with its own timing for receiving subsequent bits, without synchronization known from parallel buses. (Instead, if needed, PCIe controllers perform buffering of data incoming via various lanes, and join them into complete data afterwards. This means that, to ensure no exhaustion of those buffers, the protocol must sill involve some *rough* synchronization between the lanes, by means of dedicated control messages — that technique, however, has a much smaller impact on the overall transmission rate).

Transmission tax

Thanks to applying the serial approach, the modern buses do not require providing a separate clock signal: the timing of arriving bits is *essentially* known upfront, and potential divergences can be detected by just observing the data stream. However, such solution would cause troubles in case of a long series of bits of same value (e.g. having received 1000 "1"s in a row, we wouldn't have certainty if, by some chance, there wasn't one more or less of them). To address that, all buses discussed here employ a kind of **encoding** which guarantees periodic changes of the stream content.

The oldest variant of that, used in SATA and older version of PCIe and USB, is the **8b/10b encoding**, transforming sequences of 8 bits into 10-bit representations in such a way that the resulting stream of bits can never contain more than 5 identical values in a row. The price to be paid, however, is "wasting" as much as 20% of the total throughput. In case of SATA, the typically described maximal rate of 6 Gb/s does not take that effect into account; therefore, the actual maximal throughput of the interesting data is smaller by 20%).

An additional benefit from the 8b/10b encoding is that it acts as a **checksum**, allowing to detect a case when a single bit has been corrupted.

In newer PCIe and USB buses, we can meet newer versions of the above encoding, with a lower transmission tax, e.g. 128b/130b — using just 2 bits per every 130 bits, though requiring a better control of time on the controller side. In the PCIe 6.0 standard, yet another one encoding is used: 242B/256B — that one comes again with a higher tax, but in exchange it contains an **error correction code**, that is, a more developed system of checksums which allows not only detecting that a bit has been corrupted but also fixing problems of such kind.

Compatibility

All the abovementioned types of buses have evolved over years, introducing new versions with always growing transmission rate (altogether — measuring in Gb/s — we saw an increase from 1.5 to 6 between SATA 1 and SATA 3, from 2 to 970 between PCIe 1.0x1 and PCIe 6.0x16, and from 0.0015 to 20 between USB 1.0 and USB 3.2). The multitude of changes would lead to a nightmare for the end users, if the standards did not apply **backward compatibility** between consecutive versions. For example, a USB 2.0 device can be attached to USB 3.0 port; also conversely, a USB 3.0 can be plugged into a USB 2.0 port; in both these cases, the protocol will lead to agreeing on the older of the two versions (in our examples — 2.0), resulting in communication with acoordingly lower efficiency (yet, importantly, the connection will still work). An analogous compatibility will be usually observed between PCIe version, as well as between SATA versions. There are some exceptions, e.g. SATA 2 and SATA 3 are not fully compatible with SATA 1 (though the latter is now rarely used).

Simultaneously, the corresponding standards of **connectors** have been also evolving, with the general direction of decreasing size, but also introducing necessary adjustments for cooperating with the newer bus standards. For example, the broadly known USB-A ports were soon accompanied by so called micro USB (or more strictly: USB micro-B) plugs, which due to their smaller size became widespread in mobile devices, and then by USB-C (which, besides the new centre-symmetric and hence more convenient plug shape, also introduced additional pins that unlocked larger throughput for both data and power).

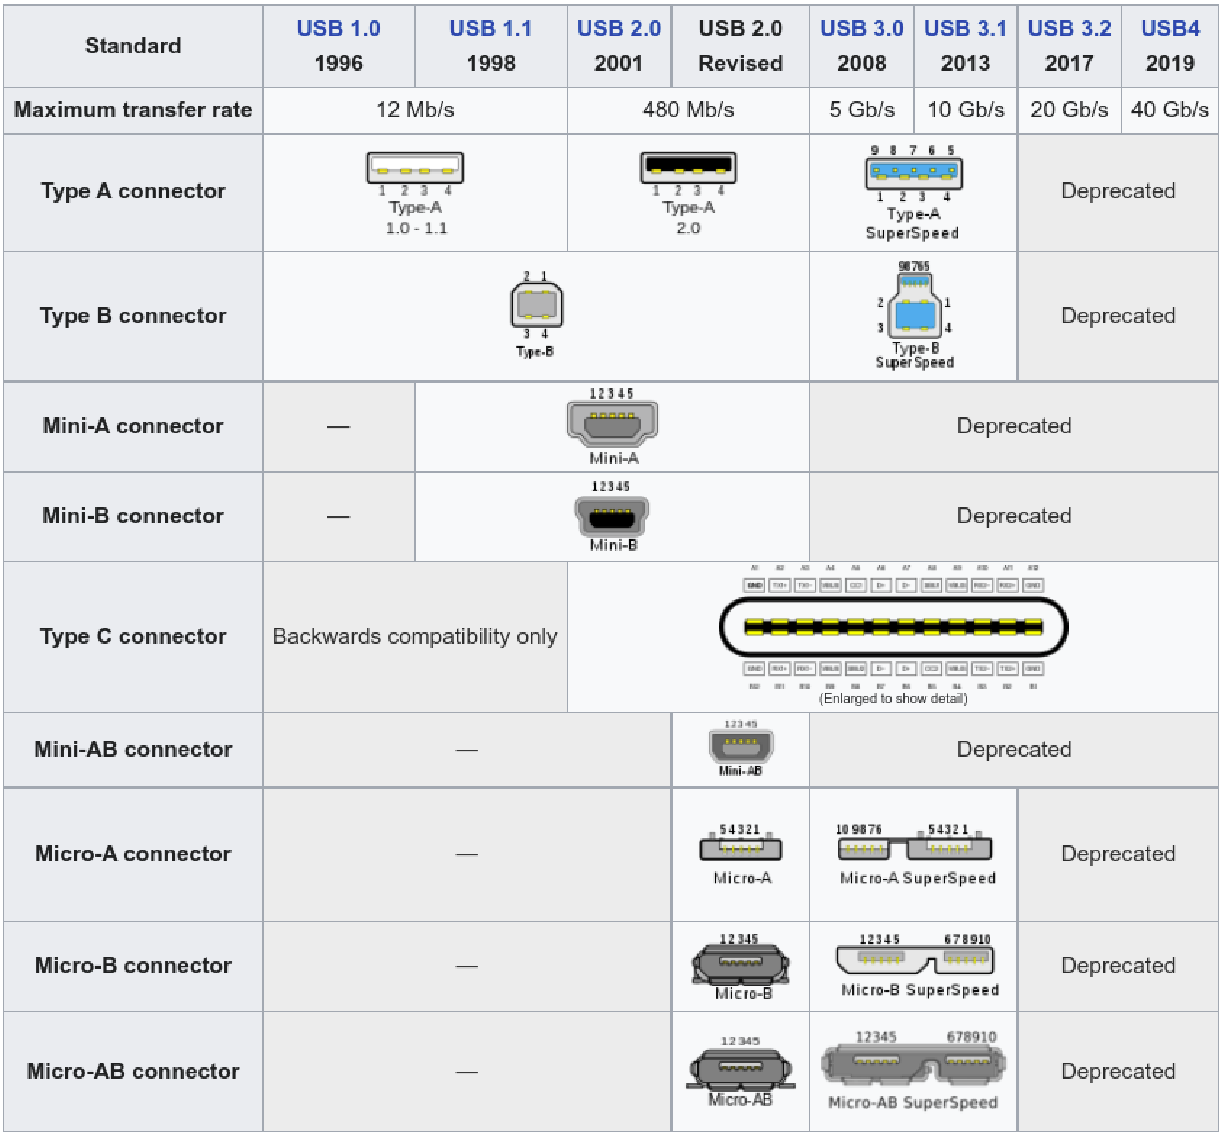

Standard	USB 1.0 1996	USB 1.1 1998	USB 2.0 2001	USB 2.0 Revised	USB 3.0 2008	USB 3.1 2013	USB 3.2 2017	USB4 2019
Maximum transfer rate	12 Mb/s		480 Mb/s		5 Gb/s	10 Gb/s	20 Gb/s	40 Gb/s
Type A connector	Type-A 1.0 - 1.1		Type-A 2.0		Type-A SuperSpeed		Deprecated	
Type B connector	Type-B				Type-B SuperSpeed		Deprecated	
Mini-A connector	—	Mini-A			Deprecated			
Mini-B connector	—	Mini-B			Deprecated			
Type C connector	Backwards compatibility only		(Enlarged to show detail)					
Mini-AB connector	—			Mini-AB	Deprecated			
Micro-A connector	—			Micro-A	Micro-A SuperSpeed		Deprecated	
Micro-B connector	—			Micro-B	Micro-B SuperSpeed		Deprecated	
Micro-AB connector	—			Micro-AB	Micro-AB SuperSpeed		Deprecated	

Figure 4. Evolution of connectors for the USB buses.
(Source: `http://en.wikipedia.org`)

Last, let us mention that the evolution of connectors can also lead to increasing compatibility, even between different bus standards. While initially SATA and PCIe were simply terminated with SATA and PCIe ports, differing in both shape and size, since 2013 there is an **M.2** standard, specifying the shape of the port (and the corresponding connector in a storage drive) together with a controller allowing communication according to any of the SATA / PCIe / USB protocols — depending on the requirements of the connected device, and potential limitations of the computer motherboard. However, while M.2 can "fit" the whole SATA 3 protocol (thus allowing its top throughput of 6 Gb/s), in case of PCIe it can only handle up to 4 transmission lanes (so for a bus of x16 type, we can only get 25% of the maximal bus throughput). Therefore, while (from the end user's viewpoint) M.2 drives can surpass the SATA barrier of efficiency, they are still behind the drives with the PCIe interface.

Logical organization of the data

From the physical viewpoint, the memory cells form just a sequence of bits (0 and 1 values). However, in case of storage drives, we also want to store them (and then interpret correctly) not just the proper data, but also: the structure of files and folders containing them, metadata of these objects, and finally additional data structures of the operating system. This means that we need an appropriate format of storing information on the drive.

File systems

Organizing data on a storage drive (e.g. choosing the location for saving new data, or encoding the relations between files and folders) is a reponisbility of the operating system, which does this following a specific set of rules, called a **file system**.

For example, for HDD drives, the Windows and Linux systems employ by default, respectively, the **NTFS** and **ext4** file systems. Until a few years ago, this meant that a HDD formatted (with default settings) and modified under Linux became non-readable under Windows; fortunately, the newest versions of Windows introduced support for ext4 (while Linux has been for long supporting NTFS, at least in read-only mode).

A disadvantage of NTFS is the relatively high fragmentation of data, which means that logically related data (e.g. from a single file) can be placed in a distant locations on the disk. As we already mentioned, that is a suboptimal situation, as it increases the total time spent on accessing a file. To address this problem, Windows offers a procedure called **defragmenting** (also called *optimizing* or in short *defragging*), which restores some order of the stored data. In the older versions of Windows, that procedure could be triggered only manually, due to how engaging it was for overall PC efficiency. Newer versions introduced an option for automatic defragmenting, though the process can be still cumbersome.

As for SSD drives, one can either use one of the leading file systems mentioned above, or choose among new file systems dedicated for SSD drives: examples include exFAT in Windows or F2FS in Linux.

Partitions

Partitions are explicit sections of the drive storage space, which — from the viewpoint of the operating system and all the running programs — are fully indepedent storage units, just as if they were physically separate drives. (For example, whenever on Windows we see several "drives" — C:, D:, E: etc. — it may happen that these are subsequent partitions of a single physical drive, just acting as different *logical drives*). In particular, each partition may be accessed with another file system.

If a partition hosts an operating system, its initial sectors are reserved for the boot loader (as we already described on Lecture 10). Similarly, if the partition has been formatted in the NTFS file system, then NTFS reserves its initial part for a file called MFT (*Master File Table*), containing metadata of every file (e.g.: its name, location on the drive, size, modification date, and attributes — such as "read-only" or "hidden"). MFT also contains a bitmap informing about free and occupied clusters on the drive. To avoid fragmentation of the MFT file, the file system reserves for it upfront a larger area (called *MFT zone*), in which it avoids placing other files. To be able to recover from a corruption of MFT content, the file system also keeps its copy (called *MFT mirror*) on the drive.

Encryption

File systems also offer an option of **encryption**, which means that the drive stores an encrypted version of the actual information. Thanks to that, an unpriviledged person will not be able to view the file contents on the drive, even if they get into the possession of our computer, unmount the drive and plug it in elsewhere, or use an alternative operating system.

A disadvantage of encryption is some slowdown of computer operation. It also shouldn't be expected that encryption will ensure "full safety". For example, even if file contents are encrypted, the structure of and names of files and folders typically are not.

It is also possible to encrypt a drive on the hardware level. One of methods for this is to use the TPM (*Trusted Platform Module*)[1], which started appearing in computers in the last decade. Every instance of TPM contains unique encryption and decryption keys, and the whole system is designed to prevent retrieving these keys from the device, as well as to prevent removing TPM from the motherboard (without destroying the hardware). Encrypting the whole drive with TPM can be done from the level of BIOS / UEFI. In this case, the booting process involves a need to decrypt the boot loader, which requires the user to provide a password and initiate storage decryption even before the main operating system has been started up. To make this happen, another operating system is introduced — a smaller one, with substantially limited functionality, and focused around security. An example of such system is BitLocker.

Generally, hardware-level encryption will be typically faster and more secure than software solutions. On the other hand, it may prevent the user from accessing their data in case when the computer is somehow damaged.

[1]Microsoft has placed TPM 2.0 among hardware requirements for MS Windows 11.

Lecture 13

On this lecture, we will look at the most important input/output devices which have not been discussed so far: graphics cards, screens, and printers. (Of course, there are also many other kinds of input/output devices which we frequently use, but we will restrict our discussion to these kinds).

Graphics cards

A **graphics card**, as the name suggests, is responsible for processing the images displayed on the screen. That is not an easy task. On Lecture 2, we described how a computer stores and operates on integers, floats and texts. Working with images also requires encoding them in a binary form. The easiest way would be to encode the color of each pixel separately (as it happens e.g. in BMP files); however, this method would be very inefficient computationally. (The importance of efficiency here is well known by every computer gamer — it is the graphics card computing power which often decides if the game will operate smoothly or *lag*). Therefore, the standard data representations used in graphics cards are more sophisticated.

Types of graphics cards

While a graphics card can be found in almost every computer, it can take one of the following two forms:

- an **integrated** graphics card: the word *integrated* represents the fact that it's a part of the processor;
- a **dedicated** graphics card: a separate module, usually exchanging data via a PCI Express bus (discussed on the previous lecture).

The advantages of an integrated card are:

- lower price;
- lower power consumption, which in case of laptops leads to:
  - smaller weight;
  - longer battery running time.

On the other hand, a dedicated graphics card offers greater computing power. It also leaves an option to be later replaced by a better model.

Hardware properties

The most important component of a graphics card is the **graphics processing unit** (**GPU**). In colloquial speech, the phrases *GPU* and *graphics card* happen to be used interchangeably. A GPU shows many resemblances to the computer CPU which we have already met, as both units play a similar role: just like CPU, GPU is responsible for performing computations. However, CPU is a general-purpose unit, fit for a variety of tasks, including those of serial nature. Every single core of a CPU can execute tasks on its own, serving as a stand-alone processing unit. On the other hand, the GPU architecture is optimized for highly **parallel** and **vector** computations, performed simultaneously on many cores. As a result, a GPU will typically contain many more cores than a CPU, often going into hundreds or even thousands.

Just like CPU, a graphics card needs to use memory. Dedicated cards are equipped with their own (separate) hardware die, whose size may currently reach even 24 GB, though typically is below 10 GB. For moderate applications, 2-4 GB should suffice. Graphics cards integrated with the CPU do not have such dedicated memory chips; instead, they are given some area of the main RAM memory by the processor, of size depending on hardware settings, and sometimes also on user preferences. For some cards, an option to control their memory size will be available through their drivers; such parameter can be also managed in the UEFI/BIOS settings.

Applications

Although the main task of graphics cards is graphics imaging, they have also found applications in other areas requiring heavy computations which admit intensive parallelization. These include implmentations of so-called *deep learning* or artificial intelligence — in particular when the data being handled are very large.

In the recent years, graphics cards have found a new important application in the so-called *mining* of crypto-currencies. Sadly, this has led to a substantial increase in graphics cards prices, and had a role in increasing the global consumption of electric power. Also, we observe an increased number of security attacks on computers, after which hackers use the overtaken machines as crypto-currency "mines".

Dedicated graphics cards also allow working with more connected monitors than 2, which is a frequent limit for integrated cards.

Screens

We will now consider the devices for displaying images, which may be called **screens**, **displays** or **monitors**. (All these terms largely overlap, though a *monitor* sounds more like a stand-alone device, not including e.g. the display/screen of a smartphone). Before describing the main technologies of manufacturing them, let us discuss some aspects common for all of them.

Pixels and the RGB space

In nearly all modern displays, the image is built from **pixels** — "points" (or, more strictly, tiny dots) of color — typically laid out in a rectangular array.

A single pixel is usually capable of showing any color of the device-supported color space — which is enabled by building it from three **sub-pixels**: a red, green and blue one. These sub-pixels are *monochromatic* (i.e. they shine with varying intensity, but in a fixed color).

That choice of sub-pixel colors results from law of optics regarding mixing colors of light (which, in turn, are based on the anatomy of human eye — which, in most individuals, employs a color recognition system based on 3 types of *cone cells*), as the following diagram shows:

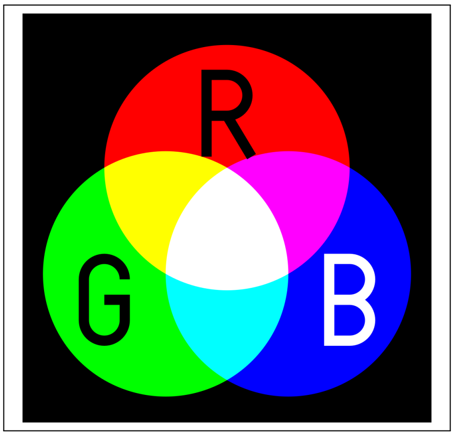

Figure 1. The diagram of mixing colors of light. (Source: `http://en.wikipedia.org`)
The above image can be obtained by placing a white screen in a fully dark room and then using three colored torches (a red, green and blue one) to illuminate three partially overlapping areas. As we see, all three primary colors combined together give the white light; combininig pairs produces yet other colors.

Generally, it turns out that *almost every* color recognizable by human eye (at least, for most people) can be obtained as a combination of lights in the 3 primary colors, just by appropriate adjustment of their intensities. This fact is the basis of the **RGB** (*Red-Green-Blue*) color system, used in screens to display various colors, and more generally, in computers to describe colors. In the most popular version, the intensity of each of the primary lights is described by an integer between 0 (none) and 255 (maximal intensity), producing a space of $256^3 \simeq 16.7$ million of possible colors. For example, the triple `(150, 30, 50)` will denote the color ■ . (Also, hexadecimal notation is commonly used here: for our example, that would be `961E32`).

One of the main screen properties from the user viewpoint is its **resolution**, which means the number of pixels displayed in the current operating mode. For example, *Full HD* resolution means a rectangular array of 1920×1080 pixels. Modern computer screens can typically operate either in their **native resolution** (the one they have been designed for) or in a *lower* one (i.e. using less pixels) — which require scaling the image and may negatively affect its quality.

Manufacturing technologies

CRT

The generation of **CRT screens** has essentially gone out of fashion. In those models, the key role is played by three electron guns emitting beams of electrons which hit a layer of *phosphor* (which means: some material which is *luminescent*, i.e. can emit light when activated by another kind of energy; not to be confused by the chemical element of *phosphorus*). The exact place where the electron beam will hit the screen is controlled by a deflecting element (using a magnetic field). Each of the electron beams periodically “runs through” all the pixels on the screen; the displayed image is stable thanks to the fact that the phosphor continues to emit light for a short time after it has been activated.

Using a **shadow mask** — a metal barrier with a grid of round holes (corresponding to pixels), placed at some distance from the screen — allows a better control of the target of the electron beams. This is even more important, given that different points of the screen are covered with phosphor shining in different RGB colors; the placement of electron guns and shadow mask holes is chosen so that e.g. the "red-light" gun (the one responsible for red sub-pixels) can hit only those points of the screen which are covered with the red-glowing phosphor.

The CRT screens were replaced by a newer technology (LCD) due to their following disadvantages:

- high power consumption;
- large size;
- eye fatigue (due to radiation);
- poor brightness and contrast;
- issues with image geometry.

LCD

The **liquid-crystal displays** (**LCD**) have no electron gun, which enabled substantially reducing their mass and changing their shape to nearly flat. Light emitted by a flat source (e.g. a system of gas-discharge CCFL lamps) falls on a polarizing filter, which lets through only light waves of specific orientation. Then, there comes a layer of liquid-crystal molecules, in which by using transistors we can turn on or off the individual pixels. (Technically, this happens as follows: by applying different voltage levels, we manipulate the pattern formed by these molecules, which impacts the polarization of light after traversing the liquid-crystal layer; that in turn determines whether this light will be blocked by yet another polarizing layer). Finally, single sub-pixels contain color filters which gives the light its desired color (red, green or blue).

There are a few technologies of managing the patterns of crystals; they influence e.g. suitable viewing angles for the screens.

There is a newer generation of liquid-crystal displays, in which the illuminating CCFL lamps are replaced with light-emitting diodes (LED). (However, their light is still processed by the liquid-crystal layer, polarizers, and often by color filters). Such devices are often called **LED displays** (though it would be more precise to call them *LED-backlit LCD displays*, to avoid confusion with "proper LED displays" in which the LED diodes directly constitute the subpixels). The advantages of the LED backlight, comparing to CCFL, are:

- better contrast (enabled by smaller size of diodes, which allows dimming them selectively in the darker areas of the image);
- better viewing angles.

The disadvantage, on the other hand, is the price, higher than for the traditional LCD displays.

Although the LCD/LED displays are usually considered better than CRT — thanks to removing the CRT disadvantages mentioned above — it's worth knowing that they also have some disadvantages on their own:

- worse representation of colors;
- worse image quality when viewing from sharp angles;

- longer response time (except best LED models);
- worse quality for non-native resolutions.

Other technologies

Let us now describe some other technologies for producing screens, which are (at least currently) less frequently used due to their price.

In the **plasma screens**, light is generated by ionizing gas, which — when brought to a plasma state — begins to emit ultraviolet light. Every subpixel is a chamber for such gas, filled with an appropriate phosphor, which upon activation by UV light begins to shine in the desired primary color.

The **OLED** screens and their improved versions (AMOLED, SUPER AMOLOED) are examples of "true LED displays", i.e. the role of subpixels is played directly by diodes of various colors. This means that we no longer need an additional "backlight" which needs to be polarized. The differences between these technologies come from different kinds of diodes. The AMOLED and SUPER AMOLED displays, due to their price, are so far used almost exclusively in small mobile devices (smartphones, tablets).

The table below shows disadvantages and advantages of plasma and OLED technologies in comparison with LCD screens.

Technology	Advantages	Disadvantages
	(in comparison with LCD)	
Plasma	• thinner • better viewing angles • better contrast • longer lifetime • better mechanical resistance	• heavier • higher power consumption • uneven usage of phosphor • image blinking • noise — noticeable on higher altitudes
OLED	• better viewing angles • better contrast (and color space) • shorter refresh time • no mercury used	• higher power consumption • more fragile • shorter lifetime • development restricted by patent rights

Ports

A screen is typically connected to a modern computer by one of the following ports:

- USB — a standard already discussed on the previous lecture.
- **HDMI** — a standard for transmitting audio & video signal (also in high quality, without compression). It has throughput up to 48 Gb/s. It also allows using longer cables (which does not just mean using more metal — but also increased chance of data corruption which needs to be detected and fixed).

- **DisplayPort** — mainly for connecting the computer to a screen or a home theater system. It allows bidirectional transfer up to 80 Gb/s. The transmitted signal may be protected from copying (with the DPCP technology) or from usage unintended by the copyright owner (with the DRM technology).

Printers

Printer colors: the CMYK space

Similarly as for screens, the printing technology is based on colored **pixels**, whose color is obtained by combining various intensities of selected primary colors. The difference is in the choice of those primary colors: since the printing inks behave more like a filter than a source of light, mixing their colors acts somewhat conversely to mixing lights, leading to *darker* results. Due to this fact, it is the standard in modern poligraphy to use the **CMY** (*Cyan-Magenta-Yellow*)) palette, which differs from the RGB set by a swap of roles between the primary and secondary colors (compare Figure 1 with Figure 2).[1]

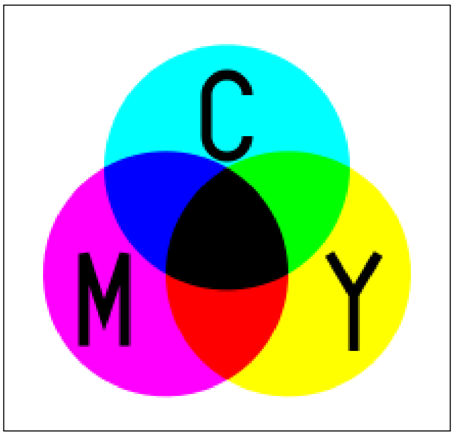

Figure 2. The diagram of mixing colors of printing ink. (Source: `http://en.wikipedia.org`)
The above image can be obtained by putting on a white sheet of paper simultaneously three circles of printing ink (a cyan, magenta and yellow one). As we see, the primary CMY colors coincide with the secondary light colors in the RGB model (see Figure 1), and conversely: combining pairs of CMY inks leads to RGB primary colors. In practice, however, this correspondence will not be perfect, due to technological differences between ink and diode producers.

In a perfect world, a just combination of cyan, magenta and yellow would be black; however, in practice, such "blackness" turns into a bit dirty brown. Moreover, a more pure black ink can be produced much cheaper than by mixing 3 colored ones. Therefore, the system is expanded by introducing a black inkc component, denoted by **K** (*key*), producing altogether the **CMYK** model. In it, each color is described with 4 numbers, specifying the intensity of the component colors, in percent (from 0 to 100). Traditionally, printing happens first with the black ink, then with magenta, cyan, and finally yellow. The important aspect here is that every ink layer should have high transparency.

[1] In the classical art theory, it is popular to consider a color system with a different set of primary colors, described now in short as RYB (*Red-Yellow-Blue*). That one has the advantage of being more intuitive (more tied to the culturally "important" palette). Technically speaking, the "quality" of such system (i.e. how well it covers the space of colors recognizable by human eye) can depend on the subtleties of choice of the primary colors, as well as on the physical properties of the inks/paints which are mixed. The CMY model is now considered optimal in this regard, and its additional advantage is its natural connection to the RGB system of light colors.

Printer types

We distinguish three main types of printers:

- **Dot matrix** printers: in these, paper moves under an ink ribbon which is pushed down by a needle to contact the paper in specific selected places. This somewhat resembles the mode of operation of all printing machines. This type of printers has fallen out of broad use, though it still has some usage niches. For example, it enables printing a few copies at once (through carbon paper), which together with its low price makes it an attractive choice for printing invoices.
- **Ink** printers: these operate by placing on the paper drops of ink, typically in primary CMYK colors. While the devices themselves are relatively cheap, their maintenance costs (i.e. the cost of ink) is high.
- **Laser** printers: these paint the paper by transferring to it a colorizing powder (**toner**) from an electrified roller. A laser beam is used to discharge specific locations of the roller, which makes them no longer attract the toner. These devices are more expensive than ink printers; on the other hand, their running costs are lower.

Connectivity

Printers may communicate with the computer e.g. via a USB cable (see Lecture 12); however, increasingly more models also support **wireless** communication. In the latter case, the following protocols may be used:

- **Bluetooth**[2] uses radio waves for transmitting data between two devices which have been *paired* together. The transmission requires little electric power, which leads to a disadvantage of relatively short communication range (up to 10 meters for the most common class of devices). The newest versions of Bluetooth standard achieve the transmission rate of up to 50 Mb/s.
- **WiFi** also uses radio waves. The main differences comparing to Bluetooth are:
  - larger range (up to ca. 90 metres);
  - higher power consumption;
  - faster transmission (though still significantly lower than for the fastest wired solutions);
  - no concept of "pairing": once a device is connected to a *local network*, it can essentially communicate with any other device in that network, though this requires an intermediate device called a **router**.

Computers and the electric power

To close the lecture, we will mention two issues related to using electrict power in computers.

[2]Fun fact: the name *Bluetooth* refers to the king Harald Bluetooth which has (for a short time) unified the Scandinavian Viking lands under his reign. Analogously, it was the intent of the Bluetooth creator (who was, by the way, Swedish) to create a protocol "unifying" various kinds of input/output devices. Even the Bluetooth symbol is a combination of Nordic runes corresponding to letters H and B (king Harald's "initials").

Failover supply

The first such issue is a threat of a **power outage**. Interrupting operation is unacceptable in case of critical infrastructure (e.g. in hospitals) but also in home conditions it may be strongly undesired or even pose a risk of damaging the hardware (e.g. a sudden power outage during operation of an SSD drive). The solution is to use a **failover** (or *uninterruptible*) **power supply**, known in short as **UPS**.

Besides keeping electricty on during an external power outage, UPS devices can also protect the hardware connected to them from other disruptions like sudden changes of network voltage. This leads to 3 main kinds of such devices:

- *Online*: separates the protected device completely from the external network. This means that the receiver is no longer threatened either by outages or voltage disruptions. However, this variant involves greater energy losses, more (acoustic) noise, and a shorter lifetime and higher price of the UPS device.
- *Offline*: the receiver device consumes electricity directly from the outer network, and UPS batteries are only charged when necessary. This means minimal energy losses and lack of noise in normal conditions. A switch to failover electricity occurs at the moment of external power outage; however, this transition has a duration of about 10 milliseconds. Also, this type does not protect against unexpected changes in voltage levels.
- *Line-interactive* and *line-interactive AVR*: improved variants of the offline solution, involving (to some extend) stabilizing the input voltage level. The reaction time in case of outage is also shorter than in offline UPS, being about 5 ms.

Cooling

Another issue related to electricity is the heat emitted during computer operation, threatening with overheating its components and damaging them. This means that stable operation of a computer requires an efficient **cooling** mechanism.

The components with largest power consumption are processors and graphics cards; that's why they are equipped with cooling systems. Typically these contain:

- a **radiator** — a comb-shaped system of metal plates which efficiently transfer heat from the processor or graphics card to the air located around these components;
- a **fan** which, by pushing the air into circulation, facilitates transmitting heat further away.

The rotation speed of the fan depends on the current temperature. In case when the fan is already rotating at its full speed, and the temperature nevertheless keeps growing, the computer (at least assuming a valid configuration) will automatically reduce its efficiency, or in extreme case switch to suspend mode or even to a full shutdown. The exact temperature thresholds can be changed in UEFI/BIOS settings — though one should remember that raising them too high brings a threat of damaging the hardware. This scheme is called **active cooling**, as opposed to **passive cooling** which involves just a radiator (without a fan). Passive cooling is practically never used in modern computers.

Yet another solution is **liquid cooling**, which involves setting up a circulation of some liquid (e.g. purified water) in a close loop, alternatingly close to the hottest computer components and much cooler masses of air. In this scheme, water takes away the heat from the computer components, and then gives it back to the air. The advantages of this approach are reduction of noise and better cooling efficiency. The disadvantages are the increased complexity and volume of the cooling system.

Lecture 14: Distributed systems

On this lecture, we will discuss an important area of technology which was not yet covered in this course, that is, networking and **distributed systems**. These allow delegating computing tasks or storing data to external machines.

Networking technology allows combining multiple computers into a single **distributed system** in which all the hardware resources can be used for more complex tasks. The word *distributed* here is intended to distinguish this case from *centralized* systems, in which all exchange of information happens via a single central computer. Yet another alternative to the distributed model is a single *supercomputer*, equipped with extremely powerful components. These approaches are presented on the following diagrams:

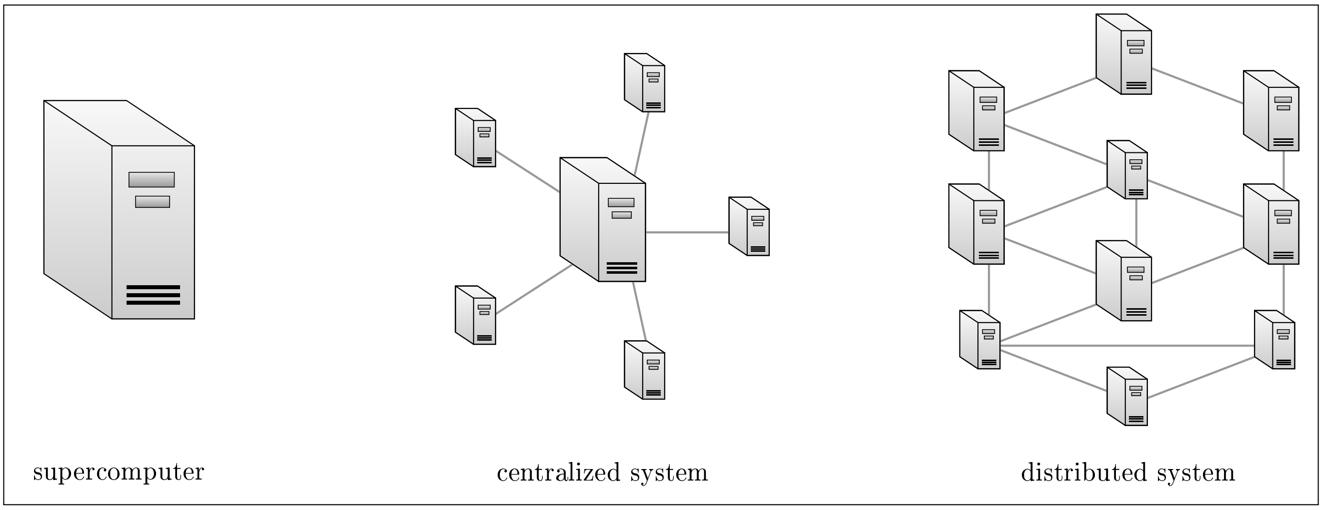

Figure 1. Types of large-scale computer systems.

Pros and cons

The advantage of distributed systems over centralized ones is typically better resilience to a failure of single node, or to network overload. In comparison to supercomputers, the advantages are following:

- increasing computing power at relatively lower costs;
- better resilience to hardware failures;
- quite easy scalability of the system
  (when more computing power is needed, just add new nodes).

On the other hand, a weakness of distrbuted systems compared to supercomputers is limited computing power of a single node, which matters in case of most complex tasks, in which parallelism may be hard to introduce, and longer latencies of transmitting data between threads may be not acceptable.

For these reasons, supercomputers are typically used for specific, most demanding tasks (e.g. complex metheorological models, intensitve training of artificial intelligence), while distributed systems are more often used for parallel execution of multiple, relatively simpler and often mutually independent, tasks.

Another disadvantage of distributed systems is the necessity for managing its complexity; moreover, in a decentralized way. This brings a whole palette of challenges, including:

- handling **asynchronous communication**, in which delivering messages takes a significant and unpredictable time (or can fail altogether);
- **replicating** information across different nodes (to make it available even in case of local hardware failures), and on the other hand, ensuring **consistency** of that information across different replicas;
- administering the system on the software level (e.g. ensuring compatibility between various operating systems, or between various versions of the same program — and/or introducing an appropriate strategy of updates).

These are all very interesting topics, though unfortunately they are outside the scope of this course.

Finally, let us note that many networks used in practice have an "intermediate" structure, that is, they are distributed systems built from *powerful* node machines (which allows efficient execution of more demanding tasks), though of course lying still far behind the top supercomputers. In such cases, when we want to expand such a system, we can choose between two strategies: **horizontal scaling** (i.e. adding new nodes, which will make the system "more distributed") or **vertical scaling** (i.e. improving the efficiency of individual nodes, which goes more in the direction of centralization).

Organzation at software level

We will now discuss a few main types of logical architecture of computer system. Depending on the ownership landscape of nodes and their role, these are:

- **Peer-to-peer architecture** (*p2p*): network nodes have diverse owners (e.g. often they will be simply personal computers of unrelated people) and have all the same role; each node may begin to be served by another node, upon establishing an appropriate bilateral conection between them.

  This model is used by popular *torrent* systems — each node in the network can send a request and start downloading data from any other node; similarly, each node can be sending files. Another example are instant messengers, like Skype.

- **Client-server architecture**: it distinguishes a central *server* (being either a single computer or — increasingly often — a whole computer system), exposing a specific functionality to other (typically numerous) nodes, which are called *clients*.

  The server is responsible for (almost) whole logic related to its functionality; in practice, the cient side may contain an additional user interface.

  For example, although the essential operation while browsing internet is sending requests and receiving responses from HTTP servers, it is also an important functionality to display the received web page in a legible way, and that is the task for a web browser (i.e. *user interface*) installed on the user's computer (i.e. on the *client side*).

In case when the server itself is a whole computer system, there are more approaches to choose from:

- **Cluster architecture**: the server consists of many machines, each of which takes its share of tasks which are similar to each other. In this model, the server itself becomes a computer network, though when speaking of a *cluster*, we typically mean a network of machines located close to each other. (In case of servers distributed to large distances, e.g. to diverse continents, we would typically describe them as a *network of local clusters*). Anyway, in all cases, it becomes important to effectively (which usually means: evenly) handle spread the work among the nodes, which is the responsibility of a *load balancer*.
- **Multi-layer architecture**: here, the functionality of the server is split into separate *modules* (or *layers*) assigned to different machines. In the simplest case, the server is split into a database and a web server communicating with the client. However, one can also distinguish a larger number of specialized layers (for example: authorization layer, cache, task scheduler).

In practice, naturally, we can meet mixed architectures, e.g. a multi-layer server in which the layer dealing with the main computations (and sometimes also other layers) will be a distributed system of a cluster architecture.

Let us also mention the **grid architecture** in which a distributed system is built of diverse machines, often belonging to various owners, which agree to dedicate some computing resources to a shared goal. Modern examples include volunteering projects for specific intensive computations (e.g. *Folding@home*, a network performing bioinformatical simulations for the purpose of fighting certain class of diseases); as another example, we can consider the block-chain networks for verifying cryptocurrencies (e.g. Bitcoin).

In the context of our classifications above, grid architectures do not allow simple labeling as a whole class; their classification depends on the particular case. Verifying cryptocurencies happens typically in a decentralized peer-to-peer network. In contrast, a typical *citizen science* project may be regarded on one hand as a single distributed *server*, whose service (computation) is used e.g. by the scientists at the root of the project; on the other hand, the individual machines involved in those computations are commonly described as *clients* which connect to a *server* (or *servers*) in order to receive subsequent shares from the central pool of tasks. (In this way — quite unusually — it is the *server* who requests work from the *client*).

Services in the cloud

In the modern distributed systems, special focus is usually put on:

- a specific **service** delivered by the system (for example: storing user's files, providing videos on demand, providing an environment for the customer's web app operation);
- ease of use, understood as hiding from the customers the whole technical layer and leaving them just with a well-defined **user interface**, through which the service is used;
- high **availability**, understood as minimizing the time of the service being inactive (due to any failures), but also minimizing the waiting time for connecting to the server and having the task done (which typically is achieved by distributing the servers over the world regions).

Such philosophy of building computer systems has a common name of **service-oriented architecture**. In the context of distributed servers of a global scope (i.e. possibly close to all users), and an implementation hidden behind an abstract interface, this technology is commonly described as the **cloud**.

Below, you can see a graph illustrating the main providers of cloud solutions. As it shows, in the case of cloud, we typically observe a single owner controlling a whole ecosystem of services and managing that whole system in a uniform way. The service delivered by the cloud provider is usually also much more extensive than in the case of grid systems.

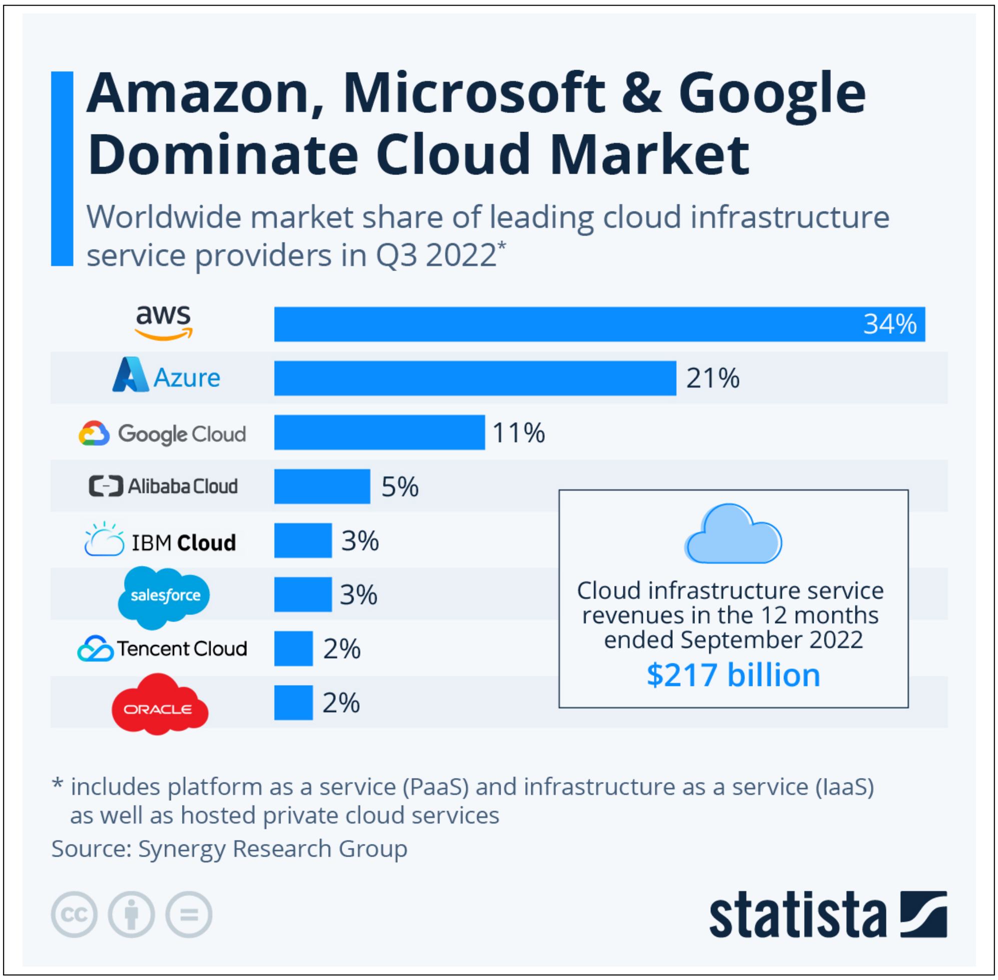

Figure 2. The leading providers of cloud services.

On the level of technical network structure, as a cloud system typically involves lots of users with independent computing tasks, cloud servers are almost always distributed systems; more specifically, clusters or distributed networks of clusters.

Levels of service

Considering the type of delivered services, we can distinguish 3 main models of cloud systems:

- **IaaS** (*infrastructure as a service*): here, the provider offers infrastructure, that is, access to servers together with maintaining them in an operational state, but no accompanying software.

  Physically, the computers could be located at the customer's headquarters, or even be customer's property. Typically, however, the customer remotely uses hardware belonging to the provider. In the latter case, it's a frequent practice to place customer's processes on *virtual* *machines* (special programs emulating a given computer model and an operating system) — which allows e.g. enforcing **isolation** between processes from different customers running on a single physical machine (as it will now suffice to run those processes on separate virtual machines).

- **PaaS** (*platform as a service*): here, in addition, the provider offers an *environment*, meaning mostly software related to the servers, network, databases. That is typically used by developers for building web applications, abstracting from technical details of the underlying hardware.

  This layer is also traditionally considered to contain utils oriented at software engineering: frameworks for automated testing, deployment, monitoring the quality of service as perceived by end user, etc.

- **SaaS** (*software as a service*): here, the provider offers specific applications used by the users for particular purposes. An example of this is Dropbox, allowing for uploading, viewing and downloading files via a web browser (thus creating a disk space "in the cloud").

As individual private users, we often use cloud solutions even without realizing that. This happens whenever we send an e-mail, listen to music from the internet, play internet games, or watch movies.

The picture becomes somewhat different when a cloud system is used by a company. Here, delegating some tasks to the cloud provider helps in areas such as:

- data analysis (regarding e.g. customer choices, application performance, etc.);
- delivering the product to end users;
- ensuring scalability of hardware
  (upwards, but also downwards — to avoid being overcharged for unused hardware resources);
- creating and maintaining data backups;
- networking security issues;
- testing the product and its new versions;
- applying the *dev-ops* culture — that is, a culture of cooperation between the developer teams and the product teams reponsible for product health and maintenance;
- applying the *agile programming* concept, which assumes e.g. frequently confronting the product with the actual users' needs, and frequently adjusting the project.

The concept of cloud is also related to the notion of **virtualization**, though we're now considering this word in a much broader sense than just *RAM virtualization* which we saw in an earlier lecture. Generally, *virtualizing* means emulating some resources by using other ones; often, by uing appropriate software. In the case of cloud solutions, this is typically implemented by the abovementioned *virtual operating system*, which ensures isolation of a single customer's processes, and may also provide a *remote desktop*, which lets the user connect to the remote server ("in the cloud") and use the emulated desktop just as if it was a part of a standard operating system. *Network virtualization* may also take place; this means splitting the capacity of physically existing network connections into independent virtual channels.